

Class 05

More on STL

fmt

Outline

- ▶ More on `std::vector`
- ▶ `std::array`
- ▶ STL algorithms
- ▶ `fmt`

std::vector

- ▶ The vector container is a STL container that provides the following:
 - ▶ Every element is adjacent in memory
 - ▶ It is resizable
 - ▶ The amortized complexity to insert something into the vector is $O(1)$
 - ▶ This is not complex at all!
 - ▶ It can manage only a single type at a time.
 - ▶ e.g. a vector cannot contain doubles, chars, and floats simultaneously
 - ▶ Size and capacity are differentiated:
 - ▶ size - how many elements are in the vector currently
 - ▶ capacity - how many elements can fit in the vector before it is full

std::vector

- ▶ There are multiple ways to create vector:

```
auto data1 = std::vector<T>{};           // empty
auto data2 = std::vector<T>{ ... };      // prepopulated
auto data3 = data1;                      // copy a vector
auto data4 = std::vector<T> data4(N);    // creates N default elements
auto data5 = std::vector<T> data5(N, V); // creates N elements equal to V
```

- ▶ Here T is a **template**. It is a placeholder for the type that the vector manages. N is an integer representing size, and V is some value of type T.
- ▶ We can create vectors that are empty, pre-filled with data, copied from another vector, set to a specific *size*, or set to a specific size with all elements equal to a specific value.

std::vector::push_back

- ▶ When adding a new element to a vector that is already filled to its capacity (e.g. using `push_back`), the vector's content must be moved in memory to a location large enough to contain the old content along with the new content. This is called *reallocation*.
- ▶ C++ will find a block of memory *twice as large, not just 1 larger*, to guarantee fewer reallocations over time. This is *always* done when a vector is automatically resized.

std::vector::reserve

- ▶ We can also tell C++ to reserve memory for a vector given some number of elements.
 - ▶ This is useful when you know how many elements you will need, but do not know what those elements are yet.
 - ▶ We use the *reserve* method to request enough memory for the specified number of elements. *This updates the vector's capacity!*

```
auto data = std::vector<double>{};  
data.reserve(100);
```

- ▶ The reserve method will find a block of memory large enough for the specified number of elements, moving any elements currently in the vector.
 - ▶ If the number of elements specified is less than the capacity of the vector, nothing happens.
 - ▶ You are still limited by your machine's hardware!

std::vector::reserve

- ▶ We can check the capacity of a vector using the *capacity* method:

```
auto data = std::vector<double>{};
data.reserve(100);
std::cout << data.capacity() << std::endl; // prints 100
```

- ▶ We can even *shrink* the vector's capacity to match the size:

```
auto data = std::vector<double>{1, 2, 3}; // capacity = 3, size = 3
data.reserve(100);                       // capacity = 100, size = 3
data.shrink_to_fit();                     // capacity = 3, size = 3
std::cout << data.capacity() << std::endl; // prints 3
```

std::vector

- ▶ What is the difference between these two vectors?

```
auto data_1 = std::vector<double>{};  
data.reserve(100);  
  
auto data_2 = std::vector<double>(100);
```


std::vector

- ▶ What is the difference between these two vectors?

```
auto data_1 = std::vector<double>{};  
data.reserve(100);  
  
auto data_2 = std::vector<double>(100);
```

- ▶ While both vectors have a capacity of 100, the first has a size of 0 and the second has a size of 100!
- ▶ The second method here creates a vector with default elements! *Note the use of parentheses, and not curly-braces*

std::vector

- ▶ When should we use one method over another?
 - ▶ Your go-to should be an empty vector, using `reserve` to preallocate memory.
 - ▶ If you plan on updating data from a baseline, then construct with default values.
- ▶ If we are going to *generate* data, then it is likely more appropriate to just create an empty vector and preallocate with `reserve`.
 - ▶ e.g. we want to generate a ballistic trajectory; we know *approximately* how many points we will need but do not know what that data will ultimately be.
- ▶ If we are going to create a set of data that we will transform and update, then it is likely more appropriate to create the vector with default values.
 - ▶ e.g. we want to implement Conway's Game of Life; we know how many points we will want *and* defaulting all cells to 0 is relevant.

Range-Based Loop

- ▶ Whenever we want to loop over some container (like a vector), we typically would iterate across the *positions* of the elements, and then get the element from its position (e.g. something like *data[i]*, where *i* is the position).
- ▶ Range-Based Loops provide a cleaner way of iterating over containers where instead of getting the positions and getting the elements from those positions, we get the element directly. They have a new syntax that we can use:

```
for (auto element_name: my_collection)
{
    // do stuff with element_name, instead of my_collection[i]
}
```

Range-Based Loop

- ▶ Example using traditional loop

```
auto data = std::vector<int>{2, 3, 5, 7, 11, 13, 17, 19};  
for (auto i = size_t{0}; i < data.size(); ++i)  
{  
    data[i] = data[i] * 2;  
}
```

Range-Based Loop

- ▶ Example using a range-based for-loop

```
auto data = std::vector<int>{2, 3, 5, 7, 11, 13, 17, 19};  
for (auto &element : data)  
{  
    element *= 2;  
}
```

Range-Based Loop

- ▶ Example using a range-based for-loop

```
auto data = std::vector<int>{2, 3, 5, 7, 11, 13, 17, 19};  
for (auto &element : data)  
{  
    element *= 2;  
}
```

- ▶ Note the ampersand (&) in the code snippet above! This is called a *reference* is necessary if you want to mutate the elements in the vector using a range-based loop.
- ▶ More on this in a later class!

How Will We Use Vectors?

- ▶ Much of what we will do in this course is generate sequences of data, usually time series data.
- ▶ Data will be buffered using vectors; this data can then be sent to other components as one large piece of data. Can be sent to:
 - ▶ algorithms (generate statistics, sorting, slicing, etc.)
 - ▶ visualizers
- ▶ Many tools deal with multidimensional data by representing the data as a series of x values, series of y values, series of z values, etc. rather than a series of points.
 - ▶ e.g. the points (1, 10), (2, 20), (3, 30) become the series (1, 2, 3) and (10, 20, 30)

STL Containers other than `std::vector`

- ▶ There are other containers in C++ that we need to mention, though we will use these a little more sparingly for various reasons.
 - ▶ `std::array`
 - ▶ `std::list`
 - ▶ `std::deque`
 - ▶ `std::forward_list`

std::array

- ▶ This is really the only other container we will care about.
- ▶ The primary differences between it and a vector are:
 - ▶ arrays are **not resizable**
 - ▶ arrays **size must be known at compile time**
- ▶ These two requirements allow arrays to outperform vector in every common use case.
- ▶ Vectors being resizable and dynamic are huge benefits, but in the scenarios where those features are not required, std::array is much better.

std::array

- ▶ To create an array, we first need to include its header:

```
#include <array>
```

- ▶ Then we can create one in a very similar way as a vector

```
auto data1 = std::array<int, 3>{};           // create array of 3 integers, all 0
auto data2 = std::array<int, 3>{10, 20, 30}; // create array of 3 integers, 10, 20, and 30
auto data3 = std::array{10, 20, 30};         // deduce type and size, not recommended
```

STL Algorithms

- ▶ The C++ STL as mentioned is packed with many algorithms implemented for you.
 - ▶ An algorithm is a set of instructions to carry out a task, whereas a function *implements* the instructions of the algorithm.
 - ▶ Many of these algorithms work with STL containers!
- ▶ Algorithms are generally found in one of the following headers in C++:
 - ▶ `#include <algorithm>`
 - ▶ `#include <cmath>`
 - ▶ `#include <functional>`
 - ▶ `#include <numeric>`
- ▶ For now, we are going to focus on algorithms meant for STL containers.

STL Algorithms

- ▶ The STL algorithm architecture for containers is unintuitive. Coders new to them expect to be able to do things like:

```
auto data = std::vector<int>{5, 7, 2, 9, 1};  
std::sort(data);    // this does not work =
```

- ▶ Unfortunately, they do not work like this. We must use special functions to retrieve the *beginning* and *end* of the container.

```
auto data = std::vector<int>{5, 7, 2, 9, 1};  
std::sort(data.begin(), data.end());    // this works!
```

- ▶ The algorithms are designed to allow them to be applied to slices of containers, rather than the entire container all the time.

STL Algorithms - Ranges

- ▶ About half of the algorithms in C++ (soon to be more!) have *ranges* variants that are much more intuitive to use. For any algorithm found in the *algorithm* header we can use a better version of the algorithm!

```
auto data = std::vector<int>{5, 7, 2, 9, 1};  
std::ranges::sort(data);
```

- ▶ Sadly in C++20 ranges only applies to the algorithm header. STL algorithms, namely found in the numeric header are missing this feature because the C++ standards committee is horribly slow.
- ▶ When possible we will use ranges over the old-style algorithms.

STL Algorithms

- ▶ STL algorithms always work in-place
- ▶ This means that when you apply an algorithm to some container, *a new container is never created*.
- ▶ For example, the following is invalid code:

```
auto data = std::vector<int>{5, 7, 2, 9, 1};  
auto data2 = std::ranges::sort(data); // Compiler error!
```

STL Algorithms Lists

- ▶ <https://en.cppreference.com/w/cpp/algorithm>
- ▶ <https://en.cppreference.com/w/cpp/numeric>

sort

- ▶ `#include <algorithm>`
- ▶ Sorts a container by placing the elements in *ascending order*.

```
auto data1 = std::vector<int>{5, 67, 2, 3, 7};  
std::ranges::sort(data1); // 2, 3, 5, 7, 67
```

- ▶ We can also sort using a predicate:

```
std::ranges::sort(data1, std::greater{}); // 67, 7, 5, 3, 2
```

- ▶ Default sorting uses the *less-than* operator to achieve *ascending order*. By overriding it to use the *greater-than* operator we can reverse sort.

min_element/max_element

- ▶ `#include <algorithm>`
- ▶ Respectively gives of the maximum and minimum elements of the container.
- ▶ Returns an *iterator* that we need to *dereference with an asterisk*.

```
auto data1 = std::vector<int>{5, 67, 2, 3, 7};  
auto min = *std::ranges::min_element(data1); // 2  
auto max = *std::ranges::max_element(data1); // 67
```

accumulate

- ▶ `#include <numeric>`
- ▶ Computes the sum of the elements in the container.
- ▶ We need to give it a starting number (usually 0!), and we need to ensure that the starting number is the correct type!

```
auto data1 = std::vector<int>{5, 67, 2, 3, 7};  
auto sum = std::accumulate(data1.begin(), data1.end(), 0); // 84
```

- ▶ Note that we are not using ranges here. Algorithms from numeric use this legacy form where we need to use begin and end positions of the container.

transform

- ▶ `#include <algorithm>`
- ▶ Applies a function (transformation) to every element in the container.
- ▶ Need to tell the transform function where to place the transformed elements.

```
auto foo(int x) -> int {  
    return x + 100;  
}  
  
auto data1 = std::vector<int>{5, 67, 2, 3, 7};  
std::ranges::transform(data1, data1.begin(), foo); // 105, 167, 102, 103, 107
```

transform

- ▶ We can use transform's output parameter to tell C++ *where* to insert the transformed elements.
- ▶ This requires one of two things to be done:
 - ▶ Initialize the destination to the appropriate size (e.g. create a vector with default elements; reserve is not enough!!)
 - ▶ Use a *back inserter*. This is a little wacky to look at, but is preferable in many cases!

```
auto data1 = std::vector<int>{5, 67, 2, 3, 7};  
auto data2 = std::vector<int>(data1.size());  
auto data3 = std::vector<int>{};  
  
std::ranges::transform(data1, data2.begin(), foo);  
std::ranges::transform(data1, std::back_inserter(data3), foo);
```

fmtlib

- ▶ *fmtlib* is an open source project that brings Python's string-formatting to C++. We will use this over cout for the rest of the semester.
- ▶ We can access fmt by including the following:

```
#include <fmt/format.h>
```

- ▶ It is much simpler and familiar to use, and we can even use it to print containers like vectors!

```
fmt::print("hello world!"); // prints "hello world"  
fmt::print("hello world!\n"); // prints "hello world" followed by a newline  
fmt::println("hello world!"); // prints "hello world" followed by a newline
```

fmtlib - parameters

- ▶ We can format strings with placeholders, and let C++ inject variables/expressions into those placeholders.
- ▶ Placeholders are specified as empty {} within the strings being formed.

```
fmt::println("{} + {} = {}", 1, 2, 1+2);

auto name = std::string{"Goku"};
fmt::println("Hello {}, how are you?", name);

auto data = std::vector<int>{1, 2, 3};
fmt::println("my vector: {}", fmt::join(data, ","));
```

fmtlib - formatting

- ▶ We can also provide special formatting for placeholders, indicating how to format the data being printed.
- ▶ Commonly we will use this to adjust the number of decimals printed with floats.

```
fmt::println("{:.4f}", std::numbers::pi); // print pi with 4 decimals
fmt::println("{:0>8d}", 9); // print 9 with 7 leading 0s
fmt::println("{:0>8d}", 9); // print 9 with 7 trailing #s
```

- ▶ The full syntax specification can be found here:
 - ▶ <https://fmt.dev/latest/syntax.html>